

B

Argus Agentic OS

Corporate Brain Design — M1

ARGUS FINANCIAL PARTNERS · CONFIDENTIAL — FOUNDER · JUNE 27, 2026

Strategic guidance only — not legal, financial, tax, or investment advice.
Securities, money-transmission, and corporate-formation decisions must be reviewed by licensed counsel and a CPA before you act. Sections flagged “see counsel” are where this matters most.

Contents

1. Mission & framing
2. Design principles
3. Human gates (CEO → CS backup)
4. Approval workflow
5. Corporate agents (C-suite brain)
6. Product squad agents
7. Agent interaction map (corporate ↔ product)
8. Decision knowledge graph (M3)
9. Model router (M4)
10. Reuse map (do NOT rebuild)
11. Milestone roadmap (CEO sign-off each deploy)
12. M1 verification checklist

Argus Agentic OS — Corporate Brain Design

Status: M1 — design + branding + webview (CEO sign-off gate)

Authoritative webview: `agentic-os/index.html`

Related: `agenticos_argus.rtf` (CEO brief)

1 Mission & framing

Mission: Run Argus Financial Partners with the smallest possible human staff — an auditable, agent-native corporation where AI agents hold real jobs, every material decision is logged, and humans gate only what must never run unattended: money out, first production launches, and final legal sign-offs.

Positioning (Phase A): Tokenization-first, non-custodial, agentic finance. We are *not* a bank — we are the all-seeing guardian of your money: keys in your device, decisions in an append-only ledger, agents that operate under scoped credentials, and a CEO who remains the final authority on financial outputs, launches, and legal posture.

Brand motif: Argus — the many-eyed sentinel. Jade accent (`#2dd4a7`) on dark surfaces; type-led hierarchy; no hype.

2 Design principles

1. **Agents decide; deterministic code executes; humans gate material outcomes** — already encoded in `runOperation` → `resolveReview` (`backend/src/operations/operationsWorkflow.ts`).
 2. **Corporate brain + product squads** — C-suite agents route portfolio work; per-product squads run PDLC.
 3. **Append-only audit everywhere** — corporate decision log + product KG; no silent overwrites.
 4. **Model-agnostic PDLC** — Claude-tiered router now; vendor seam for OpenAI/Google/local later (M4).
 5. **CEO primary, Chief of Staff backup** — three gated output classes only escalate to humans when required.
-

3 Human gates (CEO → CS backup)


Output class	Primary gate	Backup	Maps to
Financial outputs	CEO	Chief of Staff	CFO agent outputs; treasury, spend, revenue recognition
New product — first production launch	CEO	Chief of Staff	CPO + PDLC Orchestrator launch gate
Final legal signoff	CEO	Chief of Staff	CLO agent; counsel memos B4–B6

Implementation (M2): gate-policy map `(agent, output-class) → required_role` feeding `runOperation` escalation → `agent_reviews.requires_role`. Existing `SupervisionTier` values:

Tier	Meaning
<code>auto_approve</code>	Execute immediately; audit only
<code>auto_approve_audit</code>	Execute + flagged audit trail
<code>human_required</code>	Queue to <code>agent_reviews</code> ; human must approve
<code>human_led</code>	Agent drafts; human decides

4 Approval workflow

Agent invoke → gate() → escalate?



5 Corporate agents (C-suite brain)

Agent	Charter	Supervision	CEO gate
Argus Brain (Office of CEO / Orchestrator)	Routes work, convenes agents, owns CEO/CS approval queue and corporate decision log	<code>human_led</code>	Owns queue
CFO	Budgets, treasury, revenue/spend reporting	<code>human_required</code> on financial outputs	Financial outputs
CLO (General Counsel)	Legal/regulatory posture, memo drafts, filing prep	<code>human_required</code> on final signoff	Legal signoff
CISO	Corporate security posture; peers with product Cyber Specialist	<code>auto_approve_audit</code>	—
CPO	Product portfolio, roadmap, launch readiness	<code>human_required</code> on first prod launch	Launch gate
CMO	Brand, positioning, GTM (uses marketing skill)	<code>auto_approve_audit</code> (≥1k audience → admin)	—
CRO / Compliance	Risk, audit, regulatory filings	<code>human_required</code> (reuses <code>complianceSkill</code>)	—
COO / SRE	Infra, vendors, reliability	<code>auto_approve_audit</code> (reuses SRE skill)	—


Reuse today: `complianceSkill`, back-office SRE/marketing skills in `backend/src/operations/skills/`.

6 Product squad agents

Agent	Charter	Supervision
AI Product Strategist	The <i>why</i> — market, positioning, strategy docs	<code>auto_approve_audit</code>
AI Engineer	Implementation, PRs, migrations	<code>auto_approve_audit</code>
AI Spec / PDLC Orchestrator	Spec → design → build → test → launch; enforces gates	<code>human_required</code> at launch
AI Cyber Specialist	Per-product threat modeling, security review	<code>human_required</code> on material findings
AI Agentic Builder	Updates agents, skills, MCP servers continuously	<code>auto_approve_audit</code>
AI QA / Test <i>(recommended)</i>	Test plans, regression, e2e gates	<code>auto_approve_audit</code>
AI SRE / Reliability <i>(recommended)</i>	SLOs, incident summaries	<code>auto_approve_audit</code>
AI Support <i>(recommended)</i>	Support issues → product KG	<code>human_required</code> on customer-facing fixes
AI Designer / UX <i>(recommended)</i>	Quiet Premium UI, accessibility	<code>auto_approve_audit</code>

PDLC flow (M6): Strategist → Engineer + Cyber (parallel) → QA → Orchestrator launch proposal → CEO gate → production.

7 Agent interaction map (corporate ↔ product)



8 Decision knowledge graph (M3)

Append-only, in-DB (reuse append-only trigger pattern).

Nodes: Decision, Strategy, Product, Launch, Incident, SupportIssue, Fix, Agent, Approval, Filing

Edges: decided_by, rationale_for, gated_by, supersedes, relates_to, resulted_in

Write points:

Event	KG action
<code>runOperation</code> completes	<code>Decision</code> node + <code>Agent</code> edge + rationale
<code>resolveReview</code> (human gate)	<code>Approval</code> node + <code>gated_by</code> edge
Product launch (M6)	<code>Launch</code> + <code>Product</code> + <code>Strategy</code>
Support fix (M6)	<code>SupportIssue</code> → <code>Fix</code> → <code>Product</code>

Scope: Corporate KG = all decisions. Product KG = engineering, launches, strategies, support issues + fixes.

9 Model router (M4)

Registry fields: id, vendor, capability tier, context window, \$/token in+out, latency class

Routing policy: task class → required capability + token budget → cheapest qualifying model + fallback chain

Claude tiers (now):

Tier	Model class	Use when
High	Opus-class	Legal drafts, launch decisions, complex reasoning
Standard	Sonnet-class	Specs, code review, compliance analysis
Fast	Haiku-class	Triage, summaries, high-volume reads

Provider seam: `AnthropicProvider` (live) · `OpenAIProvider` (stub) · `GoogleProvider` (stub) · `LocalProvider` (stub)

Telemetry: `model_invocations` table — tokens, cost, task class → powers swap-by-usage policy.

Replaces: direct `ANTHROPIC_MODEL` calls in `skills/kycReviewSkill.ts` and peers.

10 Reuse map (do NOT rebuild)

Capability	Location
Human-gate engine	<code>operationsWorkflow.ts</code> — <code>runOperation</code> , <code>resolveReview</code> , <code>SupervisionTier</code>
Skill registry	<code>skillRegistry.ts</code> , <code>operations/skills/*</code>
Workflow engine	<code>engine.ts</code> , Temporal, Conductor
Audit	<code>auditService.ts</code> , append-only <code>audit_logs</code>
Admin agent-ops	<code>/api/admin/agent-ops</code> — extend for CEO/CS (M2)
RBAC	<code>middleware/rbac.ts</code> — add <code>ceo</code> , <code>chief_of_staff</code> (M2)
Doc render	<code>docs/build/render.mjs</code>

11 Milestone roadmap (CEO sign-off each deploy)

Milestone	Deliverable	Backend code?
M1	This doc + branding + webview	No
M2	<code>ceo</code> / <code>chief_of_staff</code> RBAC, gate-policy, Approvals admin UI	Yes
M3	<code>kg_nodes</code> / <code>kg_edges</code> , <code>decisionGraph</code> service, graph API	Yes
M4	<code>modelRouter</code> , <code>model_invocations</code> , skill refactor	Yes
M5	Corporate agent fleet as runner skills	Yes
M6	Product squad + PDLC orchestrator + product KG	Yes

12 M1 verification checklist

- ☐ Open `docs/agent-ops/index.html` in a browser (no server)
- ☐ Every corporate + product agent has a card (charter · responsibilities · interactions · tools · tier · gate)

- ☐ Three CEO-gated categories + CEO → CS backup flow visible
 - ☐ SVG org diagram + logo render; Argus branding
 - ☐ Knowledge graph + model router overviews present
 - ☐ Render `AGENTIC-OS.html` / `.pdf` via `docs/build/render.mjs`
 - ☐ CEO reviews webview + doc → signs off → M1 deployed → plan M2
-

Strategic guidance only — not legal advice. Financial and legal gates require human CEO/CS approval regardless of agent recommendation.